Schema-on-Read vs. Cloud-Native Warehousing: A Comparative Architectural Review Analysis of Apache Drill, BigQuery, and Azure Synapse

Lau Yee Wen
Faculty of Computing
Universiti Teknologi Malaysia
Johor Bahru, Malaysia
lauyeewen@graduate.utm.my

***Abstract*—The rapid growth of enterprise data has exposed the limitations of conventional batch processing pipelines, creating an urgent need for real-time analytics platforms. This paper analyzes Apache Drill, a schema-free SQL query engine, and compares its architecture and performance with cloud-native solutions, specifically Google BigQuery and Azure Synapse. Key aspects, including processing models, query latency, and usability, are evaluated to highlight the trade-offs between schema-on-read flexibility and optimized cloud data warehousing. A practical case study on cross-source data access demonstrates how Apache Drill enables seamless querying across distributed data sources without requiring data migration. The findings indicate that while Apache Drill is highly effective for ad-hoc and multi-source queries without ETL overhead, platforms such as BigQuery and Azure Synapse provide superior scalability and performance for large-scale, structured enterprise workloads. Ultimately, this comparative analysis serves as a strategic benchmarking reference for enterprise architects and IT decision-makers. It assists organizations in evaluating different analytical solutions, enabling them to make informed architectural choices that balance the need for agile, immediate data discovery with the requirements for long-term, petabyte-scale stability.**

***Index Terms*—Apache Drill, real-time analytics, Azure Synapse, Google BigQuery, cloud data warehousing**

I. Introduction

The convergence of big data and cloud computing has significantly transformed business intelligence by enabling the rapid extraction of insights from large-scale datasets [1]. However, traditional data processing architectures, which rely heavily on static ETL pipelines, struggle to handle the high-velocity data streams generated by modern Internet of Things (IoT) systems and enterprise applications [2]. As a result, there is an increasing demand for platforms capable of processing data with minimal latency [2].

Cloud-based analytics platforms address this challenge by offering scalable and high-performance solutions [3]. However, organizations must choose between highly structured data warehouses and more flexible, schema-less query engines. This paper evaluates Apache Drill as a real-time analytics solution and critically compares its architecture and performance with two leading cloud platforms: Google BigQuery and Azure Synapse.

II. Literature Review

Modern platforms for processing data in real-time and analyzing large datasets have changed significantly to handle a variety of data workloads [4].

A. Apache Drill

Apache Drill is a low-latency distributed query engine, where SQL queries can be directly performed on raw and semi-structured data (such as JSON or Parquet format) in various file systems and NoSQL databases. Unlike conventional systems, a schema-on-read architecture is adopted by Apache Drill, and the time-consuming ETL process is completely skipped.

As shown in Fig. 1, Apache Drill employs a distributed architecture consisting of daemon services called Drillbits, which are coordinated by Apache ZooKeeper. Query analysis, optimization, and execution are performed directly on raw data sources. In the figure, dotted lines represent external client query submissions, solid lines indicate internal cluster coordination via Apache ZooKeeper, and bold lines denote the high-volume extraction of raw data from storage plugins.

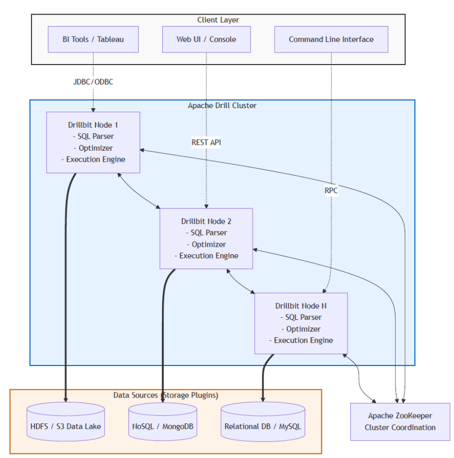

Fig. 1. APACHE DRILL ARCHITECTURE

In practical enterprise applications, interactive and ad-hoc analysis at a large scale can be achieved by systems using Apache Drill [5]. Organizations and data analysts can quickly query large, distributed datasets, like raw network logs or sensor data, without having to wait for the IT department to load the data into a central data warehouse. As a result, the bottlenecks in the data preparation process were eliminated, and the time for obtaining insights in big data operations was significantly shortened [5].

B. Cloud-Native Alternatives: Azure Synapse and BigQuery

Unlike Apache Drill's decentralised engine, a unified analysis work area has been developed by cloud service providers. For example, enterprise data warehouses are combined with big data analysis in Azure Synapse to optimise real-time data flow and processing [6]. To ensure high performance and cost-effectiveness, a specific optimisation strategy is required when such a warehouse is implemented on Azure [7]. Similarly, Google BigQuery is defined as a serverless data warehouse, where massively parallel processing (MPP) and column storage are used to achieve high-speed queries on petabyte-level data. In the latest overview of major cloud data warehouses, it is emphasised that unique architectural advantages are offered by platforms such as Google BigQuery and Azure Synapse, and a powerful, fully managed environment is provided for enterprise-scale analysis [8].

III. COMPARATIVE ANALYSIS

Although real-time analysis can be supported by Apache Drill, BigQuery, and Azure Synapse, different advantages and limitations are determined by their underlying architecture. These core differences are summarised in Table 1.

TABLE I
COMPARISON OF REAL-TIME ANALYTICS PLATFORMS

Criteria	Apache Drill	Google BigQuery	Azure Synapse
Processing Model	Schema-on-read SQL engine	Serverless data warehouse (MPP architecture)	Unified analytics workspace (Data Warehouse + Spark)
Data Ingestion	No ETL (query in-place)	Requires data loading into Google Cloud	Managed ETL/ELT pipelines
Speed & Performance	Fast, weak on big joins	Very fast for structured data	Optimized for enterprise workloads
Real-Time Streaming	Moderate (Memory-dependent, lacks built-in stream processing for heavy loads)	High via Dataflow Integration	High real-time streaming capability
Best Use Case	Quick data exploration	Large-scale analytics	Big data and BI ecosystems

The main advantage of Apache Drill lies in its schema-on-read capability, which allows newly imported JSON and semi-structured data to be queried immediately without predefined schemas. This significantly reduces data preparation time and supports rapid exploratory analysis across heterogeneous data sources. However, for production-level workloads, enterprises tend to favour cloud-based platforms such as Google BigQuery and Azure Synapse due to their optimized performance and scalability [3]. Azure Synapse provides a highly optimized environment for complex enterprise data warehousing [7], while BigQuery leverages serverless scaling to efficiently handle massive datasets without infrastructure constraints.

This difference highlights a fundamental trade-off between flexibility and performance optimization. Despite its adaptability, Apache Drill exhibits performance limitations in complex join operations. This is primarily due to its schema-on-read architecture, where query optimization occurs at runtime rather than during data ingestion. As a result, Apache Drill lacks precomputed statistics, indexing, and data partitioning strategies that are typically leveraged by cloud data warehouses. In contrast, systems such as Google BigQuery and Azure Synapse utilize cost-based optimizers and columnar storage formats, enabling more efficient execution of large-scale join queries across structured datasets.

IV. CASE STUDY AND DISCUSSION

To illustrate how Apache Drill provides real-world value, an academic study examined a large hospital system that manages billions of patient records [9]. This healthcare network needed to analyze massive amounts of complex, unstructured medical data such as doctors' text notes and sensor readings from medical devices. In traditional setups, this kind of messy data had to go through a slow and strict preparation phase, known as an ETL pipeline, where files were cleaned and organized into a central database before any research could begin.

To solve this problem, the hospital implemented Apache Drill on top of its existing Hadoop data storage systems [10]. Because Apache Drill can read raw data directly, an ability known as schema-on-read, the clinical analysts were able to bypass the time-consuming data preparation stage. They could search raw medical files immediately using standard SQL commands, which greatly reduced the time needed to discover important clinical insights without waiting for IT engineers to build strict database structures [9].

This case clearly highlights the difference between Apache Drill and cloud-native solutions. If the hospital had used Google BigQuery or Azure Synapse for the initial data search, it would have needed to organize and format all the data before uploading it into the cloud, a process known as schema-on-write. This step would have delayed the start of analysis.

However, while Apache Drill works well for fast, early data exploration, it is not designed to replace permanent, large-scale data warehouses. When a hospital needs to move from quick data searches to running very large reports such as analyzing a decade of billing records or structured global disease trends, fully managed platforms like Google BigQuery and Azure Synapse become important [6]. Google BigQuery leverages a Massively Parallel Processing (MPP) architecture, along with serverless technology, so it can divide heavy tasks across multiple computers and deliver fast, stable results [8]. In contrast, Apache Drill cannot always guarantee the same level of stability for extremely large workloads because its performance depends on the server's available memory, which can be overwhelmed during complex calculations [8].

V. CONCLUSION

The exponential growth of data volume and data processing speed makes it necessary to have an advanced analytical framework with the ability to process information in real time [4]. This comparative analysis demonstrates that Apache Drill stands out as a highly flexible, schema-less SQL engine that allows organizations to bypass conventional, time-consuming ETL processes. By allowing analysts to query raw, unstructured and distributed data in place, Apache Drill enables rapid data exploration through its schema-on-read approach, eliminating preprocessing overhead.

Although fully managed cloud-native systems such as Google BigQuery and Azure Synapse provide excellent large-scale parallel processing and optimized storage for petabyte-level enterprise workloads [6]-[8]; they still require a strict data ingestion process. Therefore, Apache Drill is more suitable for scenarios requiring rapid data exploration and schema flexibility. However, its performance limitations make it less viable for large-scale, production-critical analytics compared to cloud-native solutions. Future research should explore how schema-free engines like Apache Drill can be seamlessly integrated into the hybrid cloud data mesh architecture to combine instant data discovery with unlimited cloud computing capabilities.

REFERENCES

[1] V. Verma, "Big data and cloud databases revolutionizing business intelligence," *TIJER–International Research Journal*, vol. 9, no. 1, pp. 48–58, 2022.

[2] M. Multamäki, "Near real-time IoT data pipeline architectures," Master's thesis, 2024.

[3] N. K. Miryala and D. Gupta, "Big data analytics in cloud–Comparative study," *International Journal of Computer Trends and Technology*, vol. 71, no. 12, pp. 30–34, 2023.

[4] M. Maatallah, M. Fariss, H. Asaidi, and M. Bellouki, "Real-time data processing and big data analytics: A comparative study of modern platforms," in *The Proceedings of the International Conference on Smart City Applications*, Cham: Springer Nature Switzerland, pp. 93–101, Oct. 2024.

[5] M. Hausenblas and J. Nadeau, "Apache Drill: Interactive ad-hoc analysis at scale," *Big Data*, vol. 1, no. 2, pp. 100–104, 2013.

[6] H. P. Bomma, "Real-time data streaming and processing using Synapse Analytics," *Growth*, vol. 12, no. 6, 2024.

[7] S. T. U. Shah, "Optimizing data warehouse implementation on Azure: A comparative analysis of efficient data warehousing strategies on Azure," 2024.

[8] P. Borra, "An overview of cloud data warehouses: Amazon Redshift (AWS), Azure Synapse (Azure), and Google BigQuery (GCP)," *International Journal of Advanced Research in Computer Science*, vol. 15, no. 3, pp. 23–27, 2024.

[9] D. Chrimes, B. Moa, M. H. Kuo, and A. Kushniruk, "Operational efficiencies and simulated performance of big data analytics platform over billions of patient records of a hospital system," *Advances in Science, Technology and Engineering Systems Journal*, vol. 2, no. 1, pp. 23–41, 2017.

[10] H. Yue, "Unstructured healthcare data archiving and retrieval using Hadoop and Drill," *International Journal of Big Data and Analytics in Healthcare*, vol. 3, no. 2, pp. 28–44, 2018.

Appendix

Date	Search Keywords	AI Prompts Used	Task Description
10 April 2026	*Scopus: Apache Drill real-time analytics*	"Write me a academic manuscript title with the following information: Real-Time Analytics Platforms with Apache Drill, comparing with other existing services provided by Microsoft or Google"	Selected topic #14. Used AI to brainstorm and refine a strong, 14-word academic title that meets the rubric requirements.
11 April 2026	*Google: What is Apache Drill schema on read*	"Explain schema-on-read in simple terms."	Read basic articles to understand how Apache Drill processes data without ETL.
12 April 2026	*Google Scholar: Azure Synapse vs BigQuery*	"What is the main difference between BigQuery and Azure Synapse?"	Researched cloud data warehouses and planned this paper's structure.
13 April 2026	*IEEE Xplore: Big data analytics cloud*	"Help me outline an introduction about the growth of big data."	Drafted the Introduction section and the main problem statement.
13 April 2026	*IEEE Xplore:* Apache Drill schema-on-read	"Help me summarize the advantages of schema-on-read simply."	Found some academic papers. Wrote the literature review and comparison table.
14 April 2026	*Scopus:* Apache Drill architecture	"How does Apache Drill architecture work with ZooKeeper?"	Created the architecture diagram and wrote the technical explanation.
17 April 2026	*Google Scholar: Future of cloud data warehouse*	"What are some good points to include in a conclusion about these three tools?"	Wrote the Conclusion and finalized the comparison table.
18 April 2026	*Google:* APA 7 format guide	"Check my grammar and help me format these references to APA style."	Proofread the text, formatted the references, and finalized the IEEE layout.